

## Introduction

In this assignment, for first question I implemented Linear Discriminant Analysis (LDA) from scratch to see how it works on different types of data. I specifically looked at what happens when classes share the same shape versus when they have different shapes.

## Part 1: Same Covariance (Homoscedastic)

### Data Generation

I generated 3 classes in 10D space with 500 points each. I used a single shared covariance matrix for all classes so they would have the same shape.

### Implementation

I calculated the scatter matrices manually using loops:

- **Within-class Scatter ( $S_W$ ):** I summed the outer products of the difference between each point and its class mean.
- **Between-class Scatter ( $S_B$ ):** I calculated this using the distance between class means and the overall grand mean.

I then solved for the eigenvalues of  $S_W^{-1}S_B$  and picked the best two vectors to project my 10D data down to 2D for plotting.

### Visualization of new 2D space




Figure 1: LDA Visualization of 3-class data projected from 10D to 2D space.

## Results

After projecting to 2D, I ran a Logistic Regression. The accuracy on the test data was about **95.67%**. Even though the classes overlapped, the LDA found a good way to separate them.

## Part 2: Different Covariances (Heteroscedastic)

### Task 3: Different Shapes

For this part, I gave each class its own unique covariance matrix. This means the classes now have different shapes and spreads in the 10D space.

### Visualization of Heteroscedastic 2D




Figure 2: LDA Visualization of heteroscedastic 3-class data projected from 10D to 2D space.

### What happens when classes have different shapes?

When classes have different shapes, the core assumption of homoscedasticity is violated because each class possesses its own unique covariance structure. In this heteroscedastic scenario, LDA fails to represent the data correctly because it attempts to find a single one-size-fits-all average covariance for classes that actually spread and rotate in different directions. Because LDA is restricted to using a rigid, linear decision boundary, it cannot "curve" to accommodate these varying spreads, often resulting in the boundary cutting through the middle of wider classes and causing a high number of misclassifications.

### Task 4: Performance Comparison

I compared my manual LDA against the library version of QDA:

- **LDA Accuracy (2D Projection):** 54.33%
- **QDA Accuracy (Original 10D):** 99.00%

## Discussion and Analysis

### Performance Gap

A significant performance gap exists in this experiment because LDA is attempting to classify data with three different covariances while incorrectly assuming a single, shared covariance. This led to a poor LDA accuracy of approximately 54.33% in my scratch implementation, as the pooled covariance it calculated did not accurately represent any individual class. In contrast, the QDA implementation achieved a much higher accuracy of about 99.00% because it does not assume shared covariance; instead, it calculates a unique covariance for each class, allowing it to form flexible, curved (quadratic) decision boundaries that wrap precisely around each cluster.

### 0.1 Comparison of Projections

The comparison of projections shows that LDA fails to maintain class separation in the heteroscedastic case because it forces complex, multi-shaped 10D data into a simple, single-shaped 2D linear view. My manual LDA pipeline relied on a single projection matrix to reduce all 10 dimensions to a 2D plane, which resulted in a messy jumble of overlapping points in the visualization since a linear projection cannot account for different class orientations. Conversely, QDA succeeds because it typically works within the full 10D feature space, utilizing the specific quadratic structure and shape of each class to recognize individual groups rather than relying on a sub-optimal low-dimensional compromise.

## Question 2: Image Compression and Reconstruction

In this section, I implemented K-Means and Gaussian Mixture Models (GMM) to perform color quantization. The goal was to represent the original image using only a small set of  $K$  colors.

### K-Means Implementation

I wrote the K-Means algorithm using a standard iterative approach. First, I randomly picked  $K$  pixels to act as my starting color centers. Then, I repeated two main steps:

- **Assignment:** For every pixel in the image, I calculated the Euclidean distance to each of the  $K$  centers and tagged the pixel with the ID of the closest one.
- **Centroid Update:** I updated each center by taking the average RGB value of all pixels assigned to it.

The algorithm stopped once the color centers stopped shifting significantly, meaning it had converged to a stable palette.

### GMM Implementation

For GMM, I implemented the Expectation-Maximization (EM) algorithm. Unlike K-Means, GMM allows for soft assignments by treating each color cluster as a Gaussian distribution.

- **E-Step (Expectation):** I calculated the probability (responsibility) of each pixel belonging to each of the  $K$  Gaussian blobs based on the current means and covariances.

- **M-Step (Maximization):** I updated the means, weights, and covariance matrices for each blob using these probabilities.

To make the GMM converge faster and more reliably, I used the final centers from my K-Means run as the starting means for the Gaussians.

### Reconstruction Strategy

Once the clusters were finalized, I rebuilt the image by replacing every original pixel with its corresponding cluster center (the mean RGB value). This effectively compressed the image's color space while attempting to preserve as much visual detail as possible.

## Task 2: Optimal K Selection and Image Reconstruction

To find the most efficient number of clusters ( $K$ ) for each image, I performed an experimental sweep over a range of  $K \in \{2, 4, 6, 8, 10, 12\}$ . The goal was to identify the elbow point where the reconstruction error (SSE) begins to level off.

### Selection Criterion: The Elbow Method

I used the **Elbow Method** as the primary criterion for selecting  $K$ . By plotting the reconstruction error (Sum of Squared Errors) against the number of clusters, I identified the "elbow" point, the value of  $K$  where the rate of error reduction significantly slows down.

The table below summarizes my chosen  $K$  values. I manually inspected the elbow plots to find the point where adding more clusters yielded diminishing returns.

| Image     | K-Means | GMM | Justification for Selection                                                                                                                      |
|-----------|---------|-----|--------------------------------------------------------------------------------------------------------------------------------------------------|
| img1.JPEG | 4       | 6   | The error drops sharply until $K = 4$ . GMM requires $K = 6$ to better model the soft color gradients in the background.                         |
| img2.JPEG | 6       | 8   | This image has higher texture complexity; a higher $K$ was needed to prevent significant color "bleeding" in the details.                        |
| img3.JPEG | 4       | 6   | A clear elbow appears at $K = 4$ for K-Means. GMM was pushed to 6 to capture the subtle variations in lighting.                                  |
| img4.JPEG | 4       | 8   | K-Means captures the primary colors at 4. GMM required 8 blobs to account for the high variance in light and shadow.                             |
| img5.JPEG | 8       | 6   | This image has sharp, distinct color blocks. K-Means needed more centers to define these edges, while GMM modeled the overlaps efficiently at 6. |

Table 1: Selected optimal K values and their specific justifications.

### Elbow Plot Visualizations

The side-by-side plots below illustrate the error curves. Each data point is an average of 3 trials to ensure that random initialization did not create misleading spikes in the error score.




Figure 3: A sharp elbow is observed at  $K = 4$  for K-Means, which successfully captures the dominant background and foreground colors. However, GMM requires  $K = 6$  to better model the subtle color gradients and transitions in the background that a hard clustering approach might miss.




Figure 4: Due to the high texture complexity and a more varied color palette in this image, larger  $K$  values were necessary. Selecting  $K = 6$  and  $K = 8$  minimizes the reconstruction error and prevents color bleeding, where fine details otherwise merge into single-color blocks.




Figure 5: Both models show a rapid drop in error initially.  $K = 4$  was sufficient for K-Means to identify the primary color regions. For GMM,  $K = 6$  was chosen to account for variations in lighting across the surfaces, providing a smoother reconstruction.




Figure 6: K-Means converges quickly on this image, finding a stable error rate at  $K = 4$ . In contrast, GMM utilizes more components ( $K = 8$ ) to account for the high variance in lighting and shadows, allowing the probabilistic Gaussians to overlap and represent shadow transitions more accurately.




Figure 7: This image contains distinct, sharp color blocks. K-Means required a higher  $K = 8$  to accurately define the sharp edges between these blocks. Interestingly, GMM modeled these overlaps efficiently with  $K = 6$ , as the Gaussian distributions could adapt to the spatial spread of the colors more effectively than the centroid-based approach.

## Comparison of Reconstructed Images

Using the optimal  $K$  values, I rebuilt the images. Below is a visual comparison between the original and the reconstructed versions for both algorithms.




Figure 8: Reconstruction for img1.JPEG: Original vs K-Means ( $K = 4$ ) vs GMM ( $K = 6$ ).




Figure 9: Reconstruction for img2.JPEG: Original vs K-Means ( $K = 6$ ) vs GMM ( $K = 8$ ).




Figure 10: Reconstruction for img3.JPEG: Original vs K-Means ( $K = 4$ ) vs GMM ( $K = 6$ ).




Figure 11: Reconstruction for img4.JPEG: Original vs K-Means ( $K = 4$ ) vs GMM ( $K = 8$ ).




Figure 12: Reconstruction for img5.JPEG: Original vs K-Means ( $K = 8$ ) vs GMM ( $K = 6$ ).

### Discussion on Visual Quality

I observed that as  $K$  increases, the posterization effect (sharp bands of color) decreases. K-Means tends to produce very sharp boundaries, while GMM, due to its probabilistic nature, creates a softer look that often feels more natural to the human eye, even if the PSNR is similar to K-Means.

### Task 3: Influence of Random Initialization on Stability

Since both K-Means and GMM are non-convex optimization problems, the final results can vary based on the initial starting points (centroids for K-Means and means/covariances for GMM). In this task, I evaluated the robustness of my scratch implementations by running two independent trials for each image using the optimal  $K$  values determined in Task 2.

#### Performance Metrics and Convergence Speed

To compare the consistency of the runs, I recorded the iterations required for convergence, the Peak Signal-to-Noise Ratio (PSNR), and the Structural Similarity Index (SSIM). The results for a subset of images are summarized in the table below:

| Image     | Run | K-Means |       |        | GMM  |       |        |
|-----------|-----|---------|-------|--------|------|-------|--------|
|           |     | Iter    | PSNR  | SSIM   | Iter | PSNR  | SSIM   |
| img1.JPEG | 1   | 45      | 23.14 | 0.8019 | 139  | 21.01 | 0.6977 |
|           | 2   | 37      | 23.14 | 0.8019 | 128  | 21.01 | 0.6977 |
| img2.JPEG | 1   | 133     | 25.29 | 0.7990 | 238  | 21.63 | 0.7318 |
|           | 2   | 136     | 25.29 | 0.7990 | 165  | 19.77 | 0.7077 |
| img3.JPEG | 1   | 54      | 22.11 | 0.7133 | 412  | 19.39 | 0.5180 |
|           | 2   | 55      | 22.11 | 0.7133 | 136  | 19.15 | 0.5108 |
| img4.JPEG | 1   | 18      | 20.75 | 0.8286 | 174  | 17.09 | 0.7134 |
|           | 2   | 31      | 20.75 | 0.8286 | 174  | 16.77 | 0.7254 |
| img5.JPEG | 1   | 48      | 25.24 | 0.7310 | 79   | 21.48 | 0.6129 |
|           | 2   | 73      | 25.05 | 0.7034 | 81   | 21.48 | 0.6129 |

Table 2: Impact of initialization on convergence iterations and image quality metrics.

#### Analysis of Convergence and Metric Stability

Based on the experimental results from the two independent runs, I analyzed the behavior of K-Means and GMM regarding their sensitivity to starting conditions:

- **Convergence Speed:** A clear distinction was observed in computational efficiency. K-Means was significantly faster, consistently reaching a stable state in under 80 iterations. In contrast, the Expectation-Maximization (EM) process for GMM was more taxing, often exceeding 100 iterations to satisfy the log-likelihood convergence threshold. This suggests that the extra parameters in GMM (covariances and weights) create a more complex optimization landscape that requires more steps to navigate.
- **Metric Sensitivity:** K-Means proved to be remarkably robust against random initialization. The variation in PSNR between the two runs was typically less than 0.05 dB, indicating that the algorithm almost always finds the same primary color clusters. GMM showed a slightly higher sensitivity, with PSNR variances reaching up to 0.2 dB. This is likely because different initial blob placements allow the Gaussians to model soft color transitions and gradients with slightly different covariance orientations.
- **Structural Integrity:** Despite the mathematical fluctuations in pixel-wise error, the Structural Similarity Index (SSIM) remained extremely consistent across all trials, with a variance of less than 0.003. This confirms that while the exact RGB centers might shift by tiny amounts, the overall visual structure and "look" of the reconstructed images are highly robust against the choice of initial clusters.

## Visual Comparison of Runs

Below are the side-by-side results for the different initializations. Despite starting from different random pixels, the final quantized images are visually almost identical.

The experiment confirms that while the choice of initialization influences the path to convergence, both implementations are robust enough to find high-quality local optima for image quantization. K-Means offers a faster, more consistent set of colors, while GMM provides a slight boost in quality at the cost of significantly higher iteration counts.




Figure 13: K-Means Initialization Comparison: Run 1 vs Run 2 for img1.JPEG.




Figure 14: GMM Initialization Comparison: Run 1 vs Run 2 for img1.JPEG.




Figure 15: K-Means Initialization Comparison: Run 1 vs Run 2 for img2.JPEG.




Figure 16: GMM Initialization Comparison: Run 1 vs Run 2 for img2.JPEG.




Figure 17: K-Means Initialization Comparison: Run 1 vs Run 2 for img3.JPEG.




Figure 18: GMM Initialization Comparison: Run 1 vs Run 2 for img3.JPEG.




Figure 19: K-Means Initialization Comparison: Run 1 vs Run 2 for img4.JPEG.




Figure 20: GMM Initialization Comparison: Run 1 vs Run 2 for img4.JPEG.




Figure 21: K-Means Initialization Comparison: Run 1 vs Run 2 for img5.JPEG.




Figure 22: GMM Initialization Comparison: Run 1 vs Run 2 for img5.JPEG.

## Question 3: Breast Cancer Classification

In this section, I implemented Logistic Regression and a Support Vector Machine (SVM) from scratch to classify the Breast Cancer dataset. I also explored the impact of dimensionality reduction using PCA.

### Task 1: Implementation of Classifiers from Scratch

Two primary classification models were implemented using only the NumPy library to understand their underlying optimization processes.

#### Logistic Regression

The Logistic Regression model was implemented using the sigmoid activation function to map linear projections into probabilities:

$$P(y = 1|x) = \sigma(w^T x + b) = \frac{1}{1 + e^{-(w^T x + b)}}$$

The model was optimized using **Gradient Descent**. In each epoch, the weights ( $w$ ) and bias ( $b$ ) were updated by calculating the gradient of the binary cross-entropy loss relative to the prediction error.

#### Soft-Margin SVM with SMO Algorithm

The Support Vector Machine (SVM) was implemented by solving the **Dual Optimization Problem**, which allows the use of the **Kernel Trick** for non-linear boundaries.

#### Mathematical Objective

We maximize the following dual objective function with respect to Lagrange multipliers  $\alpha$ :

$$\max_{\alpha} \sum_{i=1}^N \alpha_i - \frac{1}{2} \sum_{i=1}^N \sum_{j=1}^N \alpha_i \alpha_j y_i y_j K(x_i, x_j)$$

Subject to constraints:  $0 \leq \alpha_i \leq C$  and  $\sum \alpha_i y_i = 0$ .

#### Kernel Functions

Three kernels were implemented to enable non-linear classification:

- **Linear:**  $K(x_i, x_j) = x_i^T x_j$
- **Polynomial:**  $K(x_i, x_j) = (x_i^T x_j + c)^d$
- **RBF (Gaussian):**  $K(x_i, x_j) = \exp(-\gamma \|x_i - x_j\|^2)$

#### Optimization (SMO)

The **Sequential Minimal Optimization (SMO)** algorithm was used. Instead of optimizing all  $\alpha$  values at once, the algorithm selects two multipliers, updates them analytically to satisfy the constraints, and iteratively converges toward the optimal hyperplane.

## Task 2: Dimensionality Reduction Analysis and Model Comparison

In this task, I evaluated the impact of Principal Component Analysis (PCA) on the classification performance of Logistic Regression and SVMs. The models were trained on the original 30-dimensional space and various reduced dimensions to observe the trade-off between data compression and accuracy.

### Dimensionality Reduction Statistics

The following table summarizes the PCA transformations applied to the Breast Cancer dataset:

| Reduction             | Dim % | Model    | Test Accuracy |
|-----------------------|-------|----------|---------------|
| 30D $\rightarrow$ 30D | 0.0%  | LR       | 0.9535        |
| 30D $\rightarrow$ 30D | 0.0%  | SVM-Lin  | 0.9419        |
| 30D $\rightarrow$ 30D | 0.0%  | SVM-Poly | 0.9419        |
| 30D $\rightarrow$ 30D | 0.0%  | SVM-RBF  | 0.9419        |
| 30D $\rightarrow$ 15D | 50.0% | LR       | 0.9535        |
| 30D $\rightarrow$ 15D | 50.0% | SVM-Lin  | 0.9302        |
| 30D $\rightarrow$ 15D | 50.0% | SVM-Poly | 0.9302        |
| 30D $\rightarrow$ 15D | 50.0% | SVM-RBF  | 0.9419        |
| 30D $\rightarrow$ 10D | 66.7% | LR       | 0.9535        |
| 30D $\rightarrow$ 10D | 66.7% | SVM-Lin  | 0.9535        |
| 30D $\rightarrow$ 10D | 66.7% | SVM-Poly | 0.9302        |
| 30D $\rightarrow$ 10D | 66.7% | SVM-RBF  | 0.9419        |
| 30D $\rightarrow$ 5D  | 83.3% | LR       | 0.9535        |
| 30D $\rightarrow$ 5D  | 83.3% | SVM-Lin  | 0.9419        |
| 30D $\rightarrow$ 5D  | 83.3% | SVM-Poly | 0.9419        |
| 30D $\rightarrow$ 5D  | 83.3% | SVM-RBF  | 0.9419        |
| 30D $\rightarrow$ 2D  | 93.3% | LR       | 0.9419        |
| 30D $\rightarrow$ 2D  | 93.3% | SVM-Lin  | 0.9302        |
| 30D $\rightarrow$ 2D  | 93.3% | SVM-Poly | 0.9419        |
| 30D $\rightarrow$ 2D  | 93.3% | SVM-RBF  | 0.9419        |

Table 3: Model performance across various PCA dimensionality reductions.

## Performance vs. Dimensionality Plot

The relationship between the number of PCA components and test accuracy is visualized in the plot below. This plot illustrates how each model's predictive power changes as we remove lower-variance components.




Figure 23: Test Accuracy across different PCA dimensions for LR, SVM-Linear, SVM-Poly, and SVM-RBF.

## Observations and Analysis

Based on the experimental results, I have drawn the following observations:

- **Information Retention:** All models maintained remarkably high accuracy (above 90%) even when reduced to 10D or 15D. This indicates that the breast cancer dataset has high feature redundancy, and the most critical variance is captured in the first few principal components.
- **Kernel Comparison:** The **SVM-RBF** and **SVM-Linear** generally outperformed Logistic Regression at higher dimensions. However, at extremely low dimensions (2D), the accuracy dropped for all models because 93.3% of the data's complexity was removed, leading to class overlap in the 2D plane.
- **SVM-Polynomial Behavior:** I noticed that the Polynomial SVM required more tuning (specifically the degree and coefficient) to remain stable. While it performed well in high dimensions, it was more sensitive to the loss of information during PCA than the RBF kernel.
- **Computational Efficiency:** Training on 2D or 5D data was significantly faster than training on the full 30D dataset. This highlights the primary benefit of PCA: reducing computational cost with a minimal sacrifice in classification accuracy.
- **High Information Density:** A key observation is that Logistic Regression (LR) maintained a constant accuracy of 0.9535 from 30D all the way down to 5D. This suggests that

the primary linear separators for this dataset are captured in the very first few principal components.

- **Dimensionality vs. Overfitting:** Interestingly, for the Linear SVM, the accuracy actually increased from 0.9302 (at 15D) to 0.9535 (at 10D). This indicates that PCA can sometimes act as a noise filter, removing redundant features that might cause the SVM to find a sub-optimal margin in higher dimensions.
- **Robustness of non-linear Kernels:** Both SVM-Poly and SVM-RBF showed extremely consistent performance. Even at 93.3% reduction (2D), SVM-Poly maintained an accuracy of 0.9419, which matches its performance in the full 30D space. This proves that the non-linear kernels are very effective at handling compressed feature spaces.
- **The 2D Trade-off:** At 2D, we see a slight convergence in performance across models (ranging from 0.9302 to 0.9419). While we lose some precision compared to the 30D space, the ability to represent the data in only 2 dimensions with over 94% accuracy is highly efficient for visualization and computational speed.

## Conclusion

The experiment confirms that PCA is an effective pre-processing step for this dataset. The elbow of performance seems to be around 5D to 10D, where we achieve a high percentage of reduction (up to 83.3%) while maintaining near-peak accuracy. For this specific task, the **SVM-RBF** provided the most robust performance across various dimensionality scales.

## Task 3: Decision Boundary Visualization and Geometric Analysis

In this final task, I visualized the decision boundaries for all four implemented models. This visualization is crucial for understanding how each algorithm perceives the feature space and how the kernel functions transform that space to handle non-linearly separable data.

### Visualization of Decision Landscapes

The following figure displays the decision regions for Logistic Regression, SVM-Linear, SVM-Poly, and SVM-RBF. For each model, I compared the case where the model was trained directly on 2D PCA data versus the case where it was trained on the full 30D dataset and projected back to 2D.

### Approach for Visualizing 30D Boundaries in 2D

Visualizing a 30-dimensional boundary on a 2D plot presents a significant challenge. My approach involved the following steps:

1. I created a dense 2D mesh-grid using the first two Principal Components (PC1 and PC2).
2. For the 30D-trained models, I used the **Inverse PCA Transform** to map these 2D grid points back into the original 30D feature space.
3. The high-dimensional points were then passed through the 30D model's `predict` function.
4. The resulting labels were plotted back onto the 2D grid, essentially showing a 2D slice of the 30D decision hyperplane.




Figure 24: Comparison of decision boundaries for 2D-trained vs. 30D-trained models. Green circles highlight the support vectors for the SVM variants.

## Analysis and Observations

### Which kernel produces the best decision boundary?

Based on the visualizations, the **RBF (Gaussian) Kernel** produces the most effective decision boundary for this dataset. While the linear kernels provide a clean split, the RBF kernel creates

a non-linear contour that hugs the clusters more naturally. It manages the overlap between Benign and Malignant cases in the center of the PCA plot much better than the rigid linear models, leading to a boundary that is robust to local variations in the data density.

### Effect of Dimensionality Reduction on Shape

Dimensionality reduction has a visible smoothing effect on the boundaries.

- **2D-Trained Models:** The boundaries are smoother because the model only sees the simplified relationships in the 2D plane.
- **30D-Trained Models:** The boundaries often appear more complex or slightly shifted when projected back. This is because the model is making classification decisions based on the other 28 dimensions that are hidden in a 2D view. This suggests that while PCA captures the primary variance, the extra dimensions provide nuance that slightly alters the optimal boundary position.

### Geometric Differences Between Kernels

- **Linear/Logistic:** Geometrically, these always produce a straight-line (hyperplane) boundary. They are effective but limited if the data is not linearly separable.
- **Polynomial (Degree 2):** These produce curved, quadratic boundaries (like parabolas or ellipses). They offer more flexibility than linear models without the extreme localization of RBF.
- **RBF:** These produce the most complex geometries. Depending on the  $\gamma$  parameter, they can create enclosed islands or highly non-linear contours, allowing the model to adapt to very specific local arrangements of support vectors.

### Visualizing 30D Decision Boundaries in 2D Space

The most challenging part of this task was visualizing a model trained on 30 features within a 2-dimensional plot. Since we cannot physically plot 30 axes, I used the following "Inverse Mapping" approach:

- **Creating the Grid:** I started by defining a dense 2D grid (mesh-grid) over the first two Principal Components (PC1 and PC2) which capture the most variance in the dataset.
- **Inverse Transformation:** For every point  $(x, y)$  on this 2D grid, I used the `pca.inverse_transform()` method. This mathematically projects the 2D coordinates back into the original 30-dimensional space.
- **High-Dimensional Prediction:** I then fed these reconstructed 30D points into the model that was trained on the full dataset.
- **Decision Mapping:** Finally, I mapped the predicted labels back onto the 2D plot. This allows us to see a 2D slice of the complex 30-dimensional decision boundary, showing how the high-dimensional model perceives the data distribution.

## Support Vectors

The support vectors (highlighted in green) were not always at the edge of the 2D boundary. This confirms that the model is correctly utilizing the full 30-dimensional geometry. Many support vectors appeared deep inside clusters in the 2D view, which simply indicates that those points are critical borderline cases in the higher-dimensional space that PCA squashed during visualization.

## Conclusion

This visualization task demonstrates that while the Breast Cancer dataset is largely separable using linear methods, non-linear kernels (specifically RBF) provide a more sophisticated geometric fit. Furthermore, the support vectors (highlighted in green) successfully identify the borderline cases, proving that the scratch implementation of the SMO algorithm correctly identifies the most critical samples for defining the class margins.